

4190.308: Computer Architecture
Midterm Exam
November 7th, 2017
Professor Jae W. Lee

Student ID #: _____

Name: _____

This is a closed book, closed notes exam.

90 Minutes

14 Pages

(+ 2 Pages for Appendices)

Total Score: 200 points

Notes:

- Please turn off all of your electronic devices (phones, tablets, notebooks, netbooks, and so on). A clock is available on the lecture screen.
- Please stay in the classroom until the end of the examination.
- You must not discuss the exam's contents with other students during the exam.
- You must not use any notes on papers, electronic devices, desks, or part of your body.

(This page is intentionally left blank. Feel free to use as you like.)

Part A: Short Answers (20 points)

Question 1 (20 points)

Please answer the following questions. You don't have to justify your answer—just write down your answer only.

Don't guess! You will get 4 points for each correct answer and lose 4 points for each wrong answer (but 0 point for no answer).

- (1) When performing multiple *integer* additions, the order of additions does not affect the final result since addition is commutative. (True/False)

ANSWER:

- (2) The starting address of a structure in C (i.e., `struct { ... }`) must be a multiple of 8. (True/False)

ANSWER:

- (3) RISC architectures generally have fewer registers than CISC architectures. (True/False)

ANSWER:

- (4) Variable `x` has a 4-byte value of `0xdeadbeef` and the address of `x` is `0x400`. Assuming a *big-endian* architecture, how is the value stored in memory? Fill the bytes at the right locations.

Address	0x3FC	0x3FD	0x3FE	0x3FF	0x400	0x401	0x402	0x403	0x404
Value									

- (5) Assuming `x` is an `int` variable in C, if `x >= 0`, the following inequality *always* holds: `-x <= 0`. (True/False)

ANSWER:

Part B: Floating-Point Numbers (20 points)

Question 2 (20 points)

To accelerate deep learning applications, Company M has introduced a new floating-point format, called *m\$-fp9*, which is a 9-bit floating-point representation based on the IEEE 754 standard. The most significant bit represents a sign bit. The next three bits are the exponent with an exponent bias of 3. The last five bits are the fraction. The rules are like those in the IEEE standard (normalized, denormalized, representation of zero, infinity, and NaN).

Sign (1 bit)	Exponent (3 bits)	Fraction (5 bit)
-----------------	----------------------	---------------------

(1) Fill in the empty boxes in the following table.

Number	Decimal Representation	Binary Representation
Positive Zero	+0.0	
Negative Zero	-0.0	
3.25_{10}	3.25 (13/4)	
0.125_{10}	0.125 (1/8)	
One	1.0	
Positive Infinity	$+\infty$	
Negative Infinity	$-\infty$	
The largest number		
The smallest positive number		

(2) Using the *denormalized* form, (a) how many non-zero values can be represented?; (b) what is the maximum denormal number?; (c) what is the smallest positive non-zero denormal number?

Part C: Human x86-64 Compiler (46 points)

Question 3 (22 points)

Alice Hacker wrote the following C code to run it on x86-64/Linux system.

```
#include <stdio.h>
int switch_func(int x, int y, int z)
{
    int res = 1;
    switch (x) {
        case 1:
            res = y * z;
            break;
        case 3:
            res += y;
        case 4:
            res -= z;
            break;
        case 6:
        case 7:
            res = y/z;
            break;
        default:
            res = 3;
    }
    return res;
}

int main() {
    int result1 = switch_func(3, 2, 1);
    int result2 = switch_func(2, 3, 4);
    printf("Welcome to CA world\n");
    printf("Result 1 = %d\n", result1);
    printf("Result 2 = %d\n", result2);
    printf("Result 3 = %d\n", switch_func(1, 5, 7));
    return 0;
}
```

(1) Fill in the blanks from the output of the program.

```
Welcome to CA world
Result 1 = _____
Result 2 = _____
Result 3 = _____
```

- (2) The assembly code for `switch_func()` is shown below. Fill in the jump table to make the program work correctly. The first entry is already provided as a reference.

```
switch_func:
    movl    $1, %eax
    cmpq    $8, %rdi
    ja      .L2
    jmp     *.L0(,%rdi,8)
.L2:
    movl    $3, %rax
    ret
.L3:
    movq    %rsi, %rax
    imulq   %rdx, %rax
    ret
.L4:
    subq    %rdx, %rax
    ret
.L7:
    addq    %rsi, %rax
    jmp     .L4
.L8:
    movq    %rsi, %rax
    cqto
    idivq   %rdx
    ret
```

Jump Table

```
.section    .rodata
    .align  8
.L0:
    .quad    .L2    # x == 0
```

Question 4 (24 points)

Ben Bitdiddle is writing an assembly code, `swap.s`, as shown below together with the original C code (`swap.c`). His code is currently incomplete as the part that swaps the values of `%rax` and `%rdx` is missing. Fill in the missing part in `swap.s` without using any temporary storage in either register or memory.

```
/* swap.c */
# include <stdio.h>

int main () {
    int x = 3, y = 1;
    printf("x = %d, y = %d\n", x, y); // x = 3, y = 1
    ... // swapping x and y (omitted)
    printf("x = %d, y = %d\n", x, y); // x = 1, y = 3
    return 0;
}
```

```
# swap.s
# x in %rax, y in %rdx
.main
    pushq %rbp
    mov  %rsp, %rbp # initiate procedure
    mov  $3, %rax
    mov  $1, %rdx
    callq 0x400450 <printf@plt>
    # Implement the swap between %rax (=x) and %rdx (=y)

    # finish the function
    callq 0x400450 <printf@plt>
    mov  $0x0, %eax
    leaveq
    ret
```

Part D: Human x86-64 De-compiler (44 points)

Question 5 (20 points)

Consider the source code below, where M and N are constants defined by `#define` (not shown).

```
int array1 [M][N];
int array2 [N][M];

int copy(int i, int j)
{
    array1[i][j] = array2[j][i];
}
```

Suppose the code above generates the following assembly code:

```
copy:
    pushq %rbp
    movq  %rsp, %rbp
    pushq %rbx
    leaq  (%rdi, %rdi, 8), %rdx
    salq  $2, %rdx
    movq  %rsi, %rax
    salq  $4, %rax
    subq  %rsi, %rax
    salq  $2, %rax
    movq  array2(%rax, %rdi, 4), %rax
    movq  %rax, array1(%rdx, %rsi, 4)
    popq  %rbx
    movq  %rbp, %rsp
    popq  %rbp
    ret
```

What are the values of M and N? Infer those values from the assembly code.

M = ____

N = ____

Question 6 (24 points)

In the following C code the definitions of both `struct node` and function `func` are incomplete.

<pre>typedef struct node { _____ x; _____ y; struct node *prev; struct node *next; struct node *root; } node_t;</pre>	<pre>void func (node_t *n) { node_t *m; m = _____; m->y /= 16; }</pre>
---	---

The `func` function is compiled into the following assembly code.

```
func:
    pushq %rbp
    movq %rsp, %rbp
    leaq 16(%rdi), %rax
    leaq 24(%rax), %rax
    shrw $4, 8(%rax)
    movq %rbp, %rsp
    popq %rbp
    ret
```

Fill in the blanks in the C code above. Note that there is a unique answer.

Part E: Procedure Calls (30 points)

Question 7 (30 points)

Here is a C function (`foo`) vulnerable to a buffer overflow attack, which is one of the most common security threats for today's computers. This function has a security hole as it does not check the length of the string. Here are some notes about the function.

- Function `gets(buf)` receives a character string from the standard input (keyboard) until the user hits the <enter> key and store it to `buf`. A null character (`'\0'`) is automatically appended at the end of the string to form a valid C string.
- Characters `'0'` through `'9'` have ASCII codes `0x30` through `0x39`.
- C strings are null-terminated (i.e., terminated by a character with value `0x00`).

```
void foo(int x){
    int a[3];
    char buf[4];
    a[0] = x;
    a[1] = 0xBFFFFFF2D;
    gets(buf);
    printf("a[0] = 0x%x, a[1] = 0x%x, buf = %s\n", a[0], a[1], buf);
}
```

```
00485d0 <foo>:
00485d0: 55                pushq   %rbp
00485d1: 48 89 e5          movq    %rsp,%rbp
00485d4: 48 83 ec 10       subq    $0x10,%rsp
00485d8: 53                pushq   %rbx
00485d9: 8b 45 08          movl    0x8(%ebp),%eax
00485dc: 89 45 f8          movl    %eax,0xffffffff4(%ebp)
00485df: c7 45 f4 f3 f2 f1 movl    $0xbffffff2d,0xffffffff8(%ebp)
00485e6: 8d 5d f0          leal    0xffffffff0(%ebp),%ebx
00485e9: 53                pushl   %ebx
00485ea: e8 b7 fe ff ff   callq   80484a4 <_init+0x54> # gets
00485ef: 53                pushq   %rbx
00485f0: 8b 45 f8          movl    0xffffffff8(%ebp),%eax
00485f3: 50                pushl   %eax
00485f4: 8b 45 f4          movl    0xffffffff4(%ebp),%eax
00485f7: 50                pushl   %eax
00485f8: 68 ec 90 04 08   pushl   $0x80490ec
00485fd: e8 94 fe ff ff   callq   8048494 <_init+0x44> # printf
0048602: 8b 5d ec          movl    0xfffffec(%ebp),%ebx
0048605: 48 89 ec          movq    %rbp,%rsp
0048608: 5d                popq    %rbp
0048609: c3                ret
004860a: 90                nop
```

- (1) Fill in the following table with the locations of the program values. Express them as decimal offsets (positive or negative) relative to register `%rbp`.

Program Value	Decimal Offset
<code>a</code>	
<code>a[2]</code>	
<code>x</code>	
<code>buf</code>	
<code>buf[3]</code>	
Saved value of register <code>%rbx</code>	

- (2) Consider the case where procedure `foo` is called with argument `x` as equal to `0xA1A2A3A4`, and we type “0123456789” when `gets` is invoked. Fill in the following table indicating whether each of the program values is corrupted or not by calling `gets`.

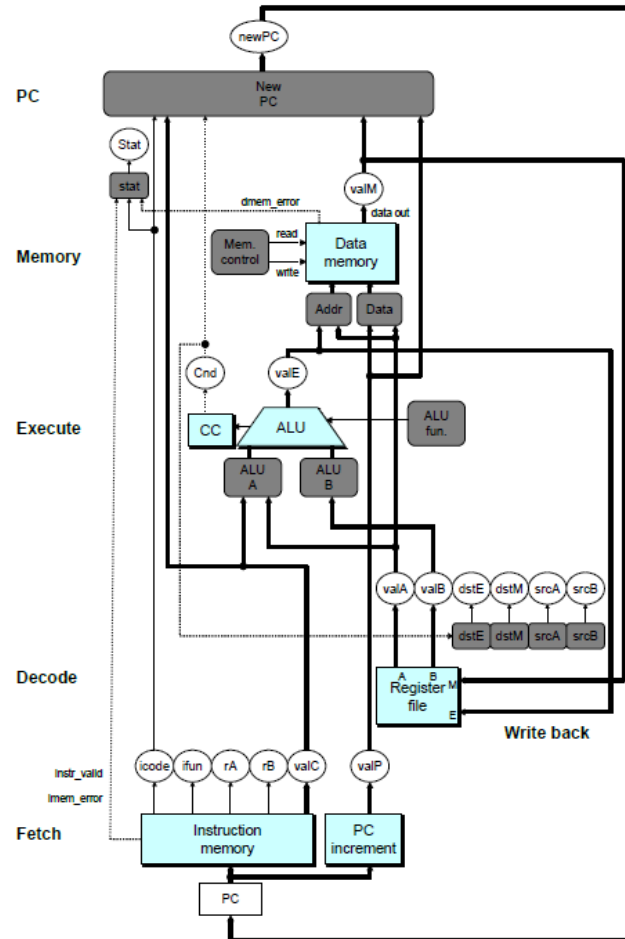
Program Value	Corrupted? (Y/N)
<code>a[0]</code>	
<code>a[1]</code>	
<code>a[2]</code>	
<code>x</code>	
Saved value of register <code>%rbp</code>	
Saved value of register <code>%rbx</code>	

- (3) Assuming the same input as in (2) what will be the program output?

- A. `a[0]` (hexadecimal): _____
 B. `a[1]` (hexadecimal): _____
 C. `buf` (ASCII): _____

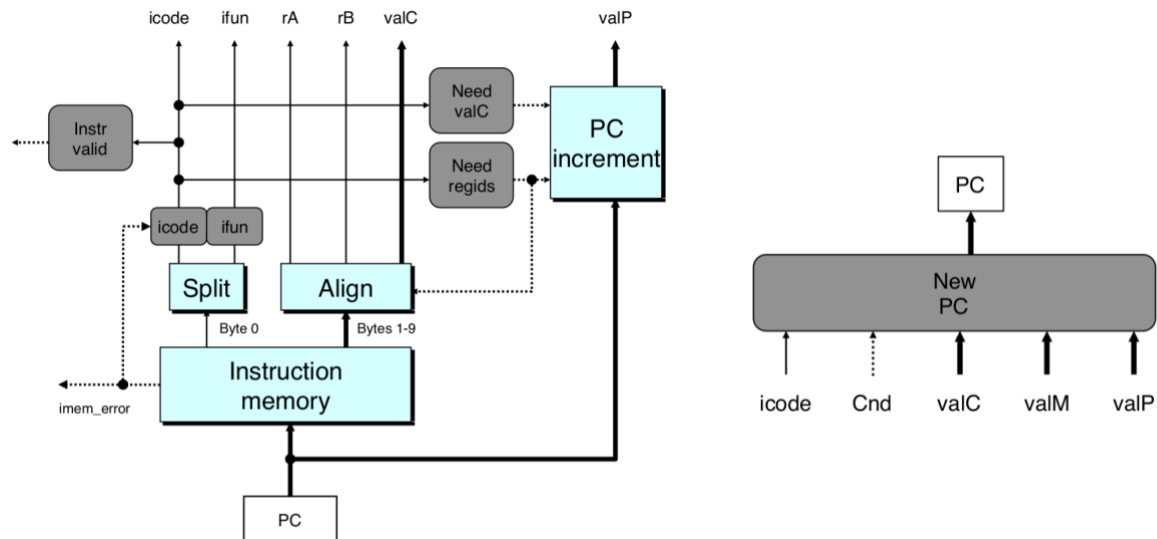
Part F: Y86-64 SEQ implementation (40 points)

The following figure shows the overall structure of Y86-64 sequential (SEQ) implementation.



Question 8 (18 points)

The following figure shows the fetch and PC update stage of the Y86-64 SEQ implementation.



- (1) Write down an HCL code for the signal `need_regids`.

```
bool need_regids =
```

- (2) Write down an HCL code for the signal `need_valC`.

```
bool need_valC =
```

- (3) Write down an HCL code for the signal `new_pc`.

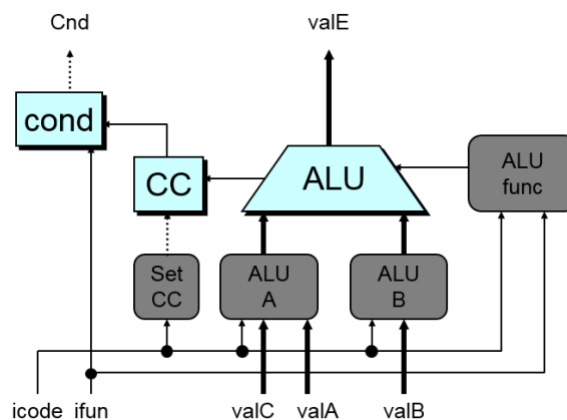
```
word new_pc =
```

Question 9 (22 points)

We would like to add `neg` instruction to the Y86-64 sequential implementation, which behaves just the same as the one in x86-64.

	icode:fn		rA:rB	
neg rA	C = INEG	0	rA	F

How should the control signals be modified in the Execute stage? Write down your code for the following four signals: SetCC, ALUA, ALUB, ALUfunc. We provide you with the original code for your reference.



	Original code	Your code
ALU A	<pre>word aluA = [icode in {IRRMOVQ, IOPQ} : valA; icode in {IIRMOVQ, IRMMOVQ, IMRMVQ} : valC; icode in {ICALL, IPUSHQ} : -8; icode in {IRET, IPOPQ} : 8;];</pre>	
ALU B	<pre>word aluB = [icode in {IRMMOVQ, IMRMVQ, IOPQ, ICALL, IPUSHQ, IRET, IPOPQ} : valB; icode in {IRRMOVQ, IIRMOVQ} : 0;];</pre>	
ALU func	<pre>word alufun = [icode == IOPQ : ifun; 1 : ALUADD;];</pre>	

Set CC	bool set_cc = icode in {IOPQ};	
-----------	--------------------------------	--

Appendix A: X86-64 assembly

Common instructions

mov	src, dst	dst = src
movsbl	src, dst	byte to int, sign-extend
movzbl	src, dst	byte to int, zero-fill
lea	addr, dst	dst = addr
add	src, dst	dst += src
sub	src, dst	dst -= src
imul	src, dst	dst *= src
neg	dst	dst = -dst (arith inverse)
sal	count, dst	dst <= count
sar	count, dst	dst >= count (arith shift)
shr	count, dst	dst >= count (logical shift)
and	src, dst	dst &= src
or	src, dst	dst = src
xor	src, dst	dst ^= src
not	dst	dst = ~dst (bitwise inverse)
cmp	a, b	b-a, set flags
test	a, b	a&b, set flags
jmp	label	jump to label (unconditional)
je	label	jump equal ZF=1
jne	label	jump not equal ZF=0
js	label	jump negative SF=1
jns	label	jump not negative SF=0
jg	label	jump > (signed) ZF=0 and SF=OF
jge	label	jump >= (signed) SF=OF
jl	label	jump < (signed) SF!=OF
jle	label	jump <= (signed) ZF=1 or SF!=OF
ja	label	jump > (unsigned) CF=0 and ZF=0
jb	label	jump < (unsigned) CF=1
push	src	add to top of stack Mem[--%rsp] = src
pop	dst	remove top from stack dst = Mem[%rsp++]
call	fn	push %rip, jmp to fn
ret		pop %rip

Instruction suffixes

b byte
w word (2 bytes)
l long/doubleword (4 bytes)
q quadword (8 bytes)

Condition flags

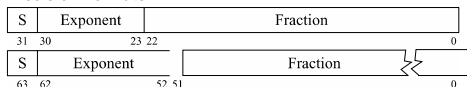
ZF Zero flag
SF Sign flag
CF Carry flag
OF Overflow flag

Suffix is elided when can be inferred from operands
 e.g. operand %rax implies q, %eax implies l, and so on

IEEE 754 FLOATING-POINT STANDARD

$$(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

where Single Precision Bias = 127,
 Double Precision Bias = 1023.

IEEE Single Precision and Double Precision Formats:**IEEE 754 Symbols**

Exponent	Fraction	Object
0	0	± 0
0	$\neq 0$	\pm Denorm
1 to MAX - 1	anything	\pm Fl. Pt. Num.
MAX	0	$\pm\infty$
MAX	$\neq 0$	NaN

S.P. MAX = 255, D.P. MAX = 2047

Registers

%rip	Instruction pointer
%rsp	Stack pointer
%rax	Return value
%rdi	1st argument
%rsi	2nd argument
%rdx	3rd argument
%rcx	4th argument
%r8	5th argument
%r9	6th argument
%r10,%r11	Caller-saved
%rbx,%rbp, %r12--%r15	Callee-saved

Addressing modes

Example source operands to **mov**

Immediate

mov \$0x5, dst

\$val

source is constant value

Register

mov %rax, dst

%R

R is register

source in %R register

Direct

mov 0x4033d0, dst

0xaddr

source read from Mem[0xaddr]

Indirect

mov (%rax), dst

(%R)

R is register

source read from Mem[%R]

Indirect displacement

mov 8(%rax), dst

D(%R)

R is register

D is displacement

source read from Mem[%R + D]

Indirect scaled-index

mov 8(%rsp, %rcx, 4), dst

D(%RB,%RI,S)

RB is register for base

RI is register for index (0 if empty)

D is displacement (0 if empty)

S is scale 1, 2, 4 or 8 (1 if empty)

source read from

Mem[%RB + D + S*%RI]

** Originally from Stanford CS107;
 modified for SNU CSE 4190.308*

Appendix B: Y86-64 (Instruction Set)

Instruction	icode:fn	rA:rB
	byte 0	1 2 3 4 5 6 7 8 9
halt	0 = IHALT	0
nop	1 = INOP	0
cmovXX rA, rB	2 = IRRMOVQ	fn
rrmovq	0	
cmovle	1	
cmovl	2	
cmove	3	
cmovne	4	
cmovge	5	
cmovg	6	
irmovq V, rB	3 = IIRMOVQ	0 F rB V
rmmovq rA, D(rB)	4 = IRMMOVQ	0 rA rB D
mrmmovq D(rB), rA	5 = IMRMOVQ	0 rA rB D
OPq rA, rB	6 = IOPQ	fn rA rB
addq	0	
subq	1	
andq	2	
xorq	3	
jXX Dest	7 = IJXX	fn Dest
jmp	0	
jle	1	
jl	2	
je	3	
jne	4	
jge	5	
jg	6	
call Dest	8 = ICALL	0 Dest
ret	9 = IRET	0
pushq rA	A = IPUSHQ	0 rA F
popq rA	B = IPOPQ	0 rA F

Register encoding

0	1	2	3	4	5	6	7
%rax	%rcx	%rdx	%rbx	%rsp	%rbp	%rsi	%rdi
8	9	A	B	C	D	E	F
%r8	%r9	%r10	%r11	%r12	%r13	%r14	No register