

Performance & Wrap-Up

Lecture 20

December 13th, 2018

Jae W. Lee (jaewlee@snu.ac.kr)

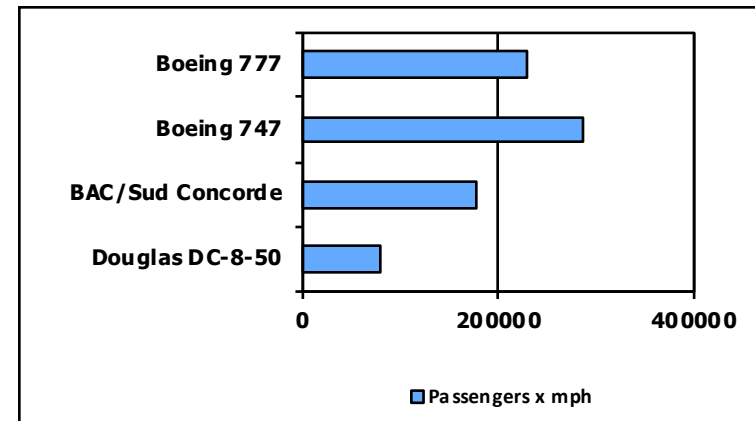



Computer Science and Engineering

Seoul National University

Slide credits: [CS:APP3e] slides from CMU; [COD5e] slides from Elsevier Inc.

Performance Example

■ Question: Which aircraft performs the best??



Today

Textbook: [P&H] 1.6

- **Performance Metrics: Time and Rate**
- **Summarizing Performance**
- **Now What?**

Performance Metrics #1: Time

- **Wall-clock time, response time, or elapsed time**
 - Actual time from start to completion
 - Includes everything: CPU time for other programs as well as for itself, I/O, operating system overheads, etc

- **CPU (execution) time**
 - CPU time spent for a given program
 - user CPU time + system CPU time
 - e.g., results of UNIX `time` command
90.7u 12.9s 2:39 65%

Performance Metrics #1: Time

■ Decomposition of CPU (Execution) Time

$$\begin{aligned}\text{CPU time} &= \frac{\text{Seconds}}{\text{Program}} \\ &= \frac{\text{Cycles}}{\text{Program}} \times \frac{\text{Seconds}}{\text{Cycle}} \\ &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}\end{aligned}$$

*This equation is called "**Iron Law of CPU Performance.**"*

Performance Metrics #1: Time

■ More on CPI (Clocks or Cycles Per Instruction)

$$\text{CPI} = \frac{\sum_{i=1}^n (\text{CPI}_i \times I_i)}{\text{Instruction Count}}$$

■ CPI Example

Instruction Class	Frequency	CPI _i
ALU operations	43%	1
Loads	21%	2
Stores	12%	2
Branches	24%	2

$$\text{CPI} = 0.43 \times 1 + 0.21 \times 2 + 0.12 \times 2 + 0.24 \times 2$$

Performance Metrics #1: Time

■ Comparing CPIs of two CPUs

- Example question: What is the CPI of CPU_S and CPU_Q?

Instruction Type	Instr. count (millions)	Cycles per Instr. (CPI)	
		CPU _S	CPU _Q
Arithmetic & Logic	10	1	1
Load & Store	5	4	2
Branch	4	2	3
Miscellaneous (7 E)	1	4	4

$$\text{CPI}_S = (10 \times 1 + 5 \times 4 + 4 \times 2 + 1 \times 4) / (10 + 5 + 4 + 1) = 2.1$$

$$\text{CPI}_Q = (10 \times 1 + 5 \times 2 + 4 \times 3 + 1 \times 4) / (10 + 5 + 4 + 1) = 1.8$$

Question: So, CPU_Q always performs better?

Performance Metrics #1: Time

■ Factors involved in the CPU Time


$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

	$\frac{\text{Instructions}}{\text{Program}}$	$\frac{\text{Cycles}}{\text{Instruction}}$	$\frac{\text{Seconds}}{\text{Cycle}}$
Program	V		
Compiler	V		
ISA	V	V	
Organization		V	V
Technology			V

Performance Metrics #1: Time

■ RISC vs. CISC arguments

■ MIPS (typical RISC) vs. VAX8700 (typical CISC)



Source : Hennessy & Patterson *Computer Architecture: A Quantitative Approach, 5th Ed.(Appencix L)*, Morgan Kaufmann, 2012

Performance Metrics #2: Rate

■ MIPS (million instructions per second)

- $\text{MIPS} = \frac{\text{Instruction count}}{\text{Execution time} \times 10^6}$

$$\text{Execution time} \times 10^6$$

- Specifies performance (roughly) inversely to execution time
- Easy to understand; faster machines means bigger MIPS
- Problems
 - It does not take into account the capabilities of the instructions.
 - It varies between programs on the same computer.
 - It can even vary inversely with performance!!

■ MFLOPS (million floating-point operations per second)

Performance Metrics: Ratio

- “X is n times faster than Y” means:

$$\frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = n$$

- “X is n% faster than Y” means:

$$\frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = 1 + \frac{n}{100}$$

- “X is n order of magnitude faster than Y” means:

$$\frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = 10^n$$

Summarizing Performance

- Arithmetic mean
(Time)

$$\frac{1}{n} \sum_{i=1}^n T_i$$

- Harmonic mean
(Rate)

$$\frac{n}{\sum_{i=1}^n \frac{1}{R_i}}$$

- Geometric mean
(Ratio)

$$\sqrt[n]{\prod_{i=1}^n \text{Ratio}_i}$$

Summarizing Performance: Arithmetic Mean

■ Used to summarize performance given in times

- Average Execution Time = $(\sum_{i=1}^n \text{Execution Times}) / n$
- Assumes each benchmark is run an equal no. of times

■ Weighted Arithmetic Mean

- Weighted Average Execution Time = $\sum_{i=1}^n (W_i \times \text{Execution Times}) / \sum_{i=1}^n W_i$
- One possible weight assignment: equal execution time on some machine

Summarizing Performance: Harmonic Mean

■ Used to summarize performance in rates (e.g., MIPS, FLOPS):

- Harmonic Mean = $n / \sum_{i=1}^n (1 / R_i)$
- Example
 - Four programs execute at 10, 100, 50 and 20 MFLOPS, respectively
 - Harmonic mean is $4 / (1/10 + 1/100 + 1/50 + 1/20) = 22.2$ MFLOPS

■ Weighted Harmonic Mean

- Weighted Harmonic Mean = $\sum_{i=1}^n W_i / \sum_{i=1}^n (W_i / R_i)$

Summary: Performance

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

- **“Execution time is the only and unimpeachable measure of performance”**
 - CPU time equation can predict performance by estimating the effects of changing features.
- **Measuring performance requires good care**
 - Good ways to summarize performance
 - Good workloads (benchmarks)

Today

Textbook: [P&H] 1.6

- Performance Metrics: Time and Rate
- Summarizing Performance
- **Now What?**

In Lecture 1...

■ Do you remember this one?

Seoul National University

Why you should take this course?

- Because....You won't graduate if you don't take this course.
- Because....You want to design the next great instruction set.
 - Instruction set architecture has largely converged, especially in the desktop/server/laptop space.
 - Dictated by powerful market forces (Intel/ARM).
- Because....You want to become a computer architect and design the next great computer systems.
- Because....The design, analysis, implementation concepts that you will learn are vital to all aspects of computer science and engineering – operating systems, computer networks, compiler, programming languages
- Because....The course will equip you with an intellectual toolbox for dealing with a host of systems design challenges
- **And much more !!!**

Source: Prof. Fernando C. Colon Osorio's lecture notes

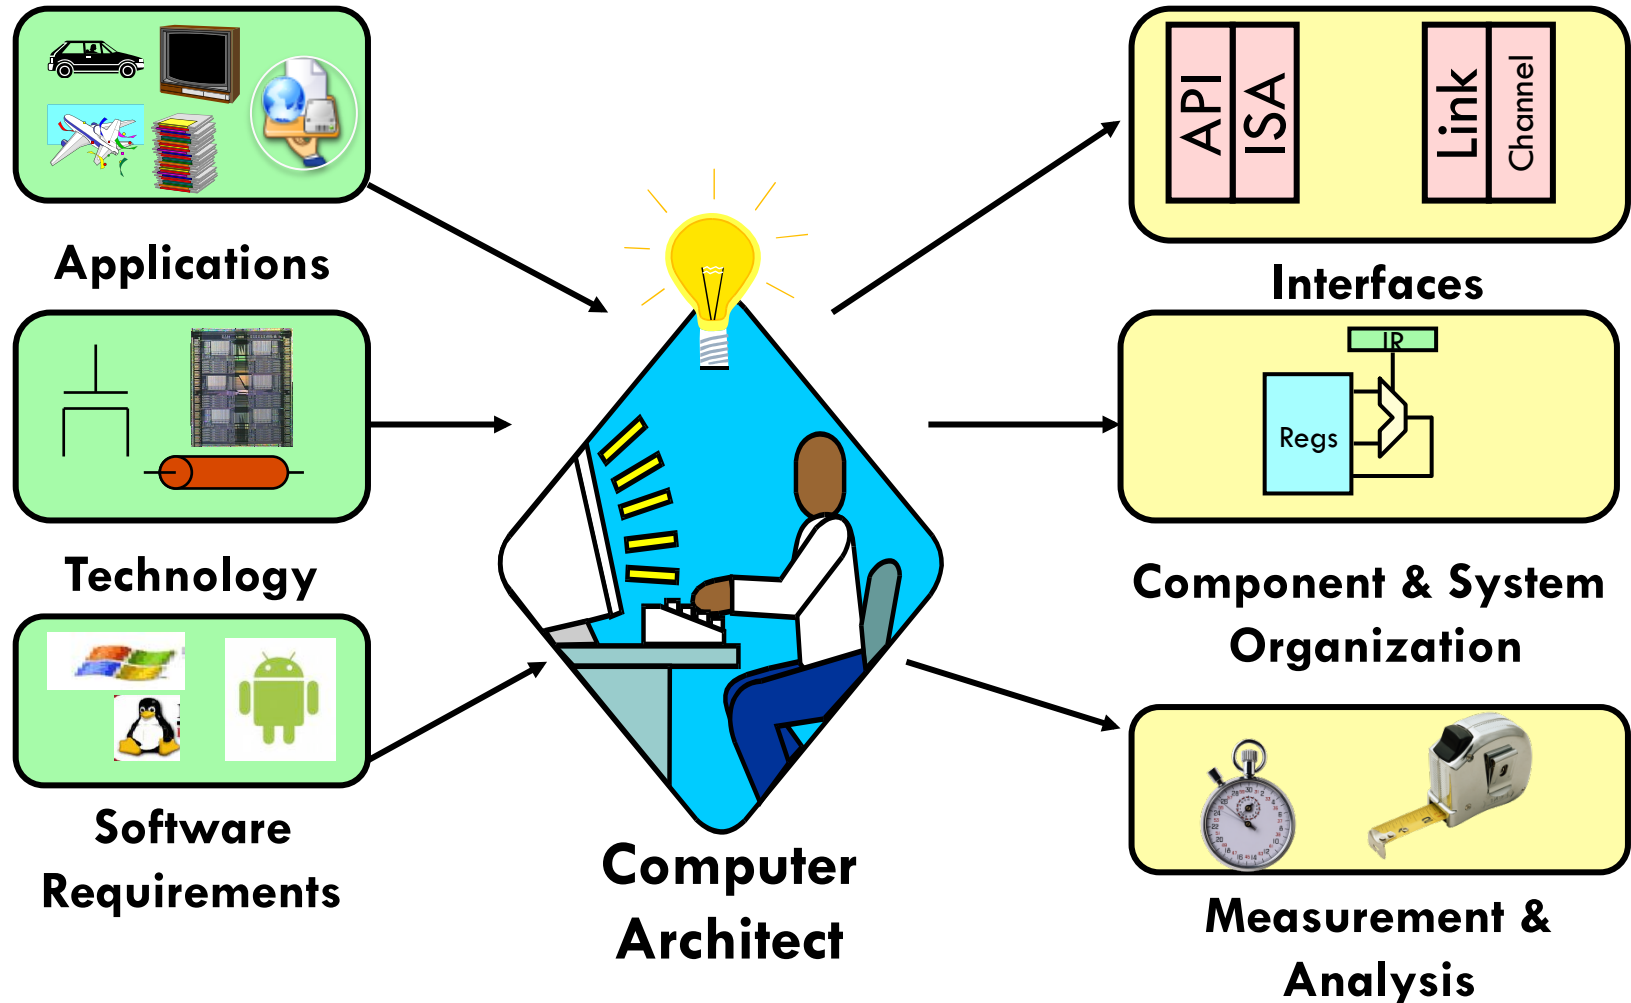
Where Are We Now?

■ Here is what we learned:

- Instruction Set Architecture (i.e., abstraction of hardware)
- Representing numbers: integer and floating-point
- x86-64 assembly and how to translate a C program into it
- Basic processor organization
- Branch prediction
- Pipelining
- Locality and memory hierarchy
- Caches
- Program optimization for caches
- Virtual memory
- Performance evaluation

■ I hope those tools provide a solid foundation for your CSE-related careers (as they did for mine).

What Computer Architects Do?



Source: Stanford EE282 (Prof. C. Kozyrakis)

겨울방학 인턴/UROP 모집 공고

- **아키텍처 및 코드 최적화 연구실: 약간명 (1-2명)**
 - 빅데이터 프로세싱 (Apache Spark) 아키텍처 가속 기술 연구
 - JavaScript 기반 IoT 응용을 위한 CPU 아키텍처 연구
 - RISC-V CPU+Deep Learning Accelerator의 FPGA 프로토타이핑

- **학부행정실 UROP 공고 참고 (또는 contact professor)**

One More Thing

■ Final Exam

- Date: 12/18 (Tue) 10 AM – noon (2 hours)
- Place: This classroom (#302-208)
- Scope: Everything not in the scope of the midterm
- Sample finals (with solutions) provided to help your study.

■ Please fill in course evaluation.

- I would appreciate your constructive feedback for the course.
- I promise to LISTEN and IMPROVE the course based on it.

■ It was a great pleasure to teach you and I wish best of luck for all of your future endeavors! 😊